

Audio Direction – Dynamic Matching to Content

The differentiator. Audio is **not** a flat mood bed chosen for the whole clip. Music follows the narrative arc, ambient tracks what's on screen scene-by-scene, and accents hit the pop-ups. It's matched **by construction**, because the same Claude plan that writes the script and the scene visuals also writes the **audio map** against the same timeline.

Both layers are optional, set in the "Add Sounds" step. Music is off/on; **SFX/ambient is off/on and off by default** — it's a deliberate choice the user makes, never a standard insert. A video can ship with VO only, VO + music, or the full scored mix. When a layer is off, its map entries aren't generated and it costs nothing.

The audio map (emitted by `script.generate`)

Alongside the script, scenes, and captions, the planning stage emits:

```
{
  "music": {
    "style": "<from client sonic identity>",
    "sections": [
      { "label": "hook",
        "start_s": 0, "end_s": 3, "intent":
        "tension, sparse, pulse" },
      { "label": "insight",
        "start_s": 3, "end_s": 18, "intent":
        "steady bed, low energy, let VO lead"
      },
      { "label": "example",
        "start_s": 18, "end_s": 32, "intent":
        "warmer, human, subtle melody" },
      { "label": "resolution",
        "start_s": 32, "end_s": 40, "intent":
        "lift, build" },
      { "label": "closing",
        "start_s": 40, "end_s": 45, "intent":
        "resolve, leave a question" }
    ]
  }
}
```

```
    },  
    "ambient": [  
      { "scene": 1, "start_s": 0,  
        "end_s": 12, "bed": "open-plan office  
murmur, distant keyboards" },  
      { "scene": 2, "start_s": 12,  
        "end_s": 28, "bed": "quiet boardroom  
room tone" }  
    ],  
    "accents": [  
      { "at_s": 22.0, "cue": "soft  
whoosh", "trigger": "popup_in" }  
    ]  
  }  
}
```

Every entry is derived from content the model already wrote — the music intent from the narrative section, the ambient bed from that scene's visual description, the accents from the pop-up cues. Nothing is random.

1. Music that follows the arc

Generate **one coherent track** with ElevenLabs Music v2, prompted section-by-section from `music.sections`. Music v2 supports long-form section composition and mid-track

transitions,
so intensity shifts with the spine and the
**section boundaries are placed at the script's
beat timestamps** — the build into "resolution"
lands where the resolution actually starts.

- **Sonic identity.** The `style` comes from the client's brand sonic identity (a Music finetune or a locked style descriptor), so all of a client's videos share a sound. Per-video arc + per-client identity = coherent body of work, not stock-music roulette.

2. Ambient matched to each scene

For each `ambient` entry, generate the bed with the ElevenLabs SFX endpoint from the scene's visual description, sized to the scene's duration (SFX can be generated and looped for atmosphere). Beds **crossfade at scene boundaries** so the soundscape changes when the picture changes — trading floor to boardroom, street to interior.

3. Sync, accents & the mix (at assemble)

The Shotstack timeline binds it together against real timestamps:

- **Duck music under VO.** Music sits ~12–18 dB below narration during speech via level automation (sidechain-style), rising in the gaps and on the closing.
- **Ambient under music** at a low, constant bed level; crossfaded per scene.
- **Accents on pop-ups.** Each accents cue fires a short SFX at the pop-up's entrance — audio and motion graphic hit together.
- **Cut alignment.** Music section changes and ambient crossfades are placed on scene cuts, using the scene timestamps + VO word-level timestamps from the audio stage.

Why this is matched, not guessed

The chain is: content → plan (script + visuals + audio map, one model, one timeline) → generate each audio element from its map entry → render to the exact timestamps. The audio can't drift from the content because it was authored against the same beats. That's the difference between "has background music" and "scored."

Cost note

Music and ambient are duration-priced; both are dials (off / on) in the brief and estimate. Generate audio once at plan-approved length; don't regenerate for the draft-vs-final split (only the visual render changes resolution). See `cost-controls.md`.
EOF