

GAS Influencer Intelligence Platform — Build Brief

Entry point for Claude Code. Read this first, then the docs in `/docs` for the stage you're building. Build in the phase order at the bottom — do not skip ahead.

Mission

An agency-grade studio that (1) builds **believable, human** AI influencers, (2) learns a client's knowledge, and (3) **autonomously produces** publish-ready 45-second videos — a-roll (avatar), b-roll (cinematic), optional music, optional SFX/ambient, optional glass-box pop-ups — with a guiding **producer co-pilot**, a **hard cost gate** before every production, and a **world-class build UX**.

Built first for GAS Marketing's own team at `influencers.gasmarketing.co.za`, running our clients' work. **Iteration 2** opens it to the market and adds the performance/learning loop, Meta feedback, and ad-loading (this is NOT a Meta-integrated platform yet — see Phase 2).

"World's best" non-negotiables for Iteration 1:

- **Believable humans** — embedded hyper-realism prompts (skin texture, pores, natural no-makeup look) so influencers don't read as AI. See `influencer-builder.md`.
- **Audio matched to content, by construction** — music follows the arc; SFX/ambient is optional and per-scene. See `audio-direction.md`.
- **A producer in the room** — a co-pilot that helps with scene layout, camera angles, shots.
- **World-class UI/UX** — see `ux-flow.md`.
Cost control is dials, never a quality cap.

Locked stack (do not relitigate)

Concern	Choice
App + API + runtime	Next.js on Vercel (deploys from github.com on push)
Relational DB + vectors	Neon Postgres (Vercel Marketplace) + pgvector
File storage	Vercel Blob (R2 is the later egress escape hatch)
Orchestration	Inngest (durable steps on Vercel)
Auth	Auth.js (NextAuth)
Producer + script + audio-map brain	Claude Sonnet 4.6 default (Opus 4.8 premium toggle; Haiku 4.5 for ingestion) + prompt caching
Voice / music / ambient	ElevenLabs (TTS + Music v2 + SFX)
Avatar (a-roll) + photo avatar	HeyGen Avatar V API

Concern	Choice
Influencer identity + b-roll	Higgsfield Soul 2.0 (identity) + Seedance 2.0 / DoP (b-roll) + Cinema Studio (camera) – via API, NOT the Supercomputer agent
Realism enhancement (selective)	Magnific / Freepik API (Skin Enhancer + Precision upscale on hero frames)
Stitch / captions / mix / overlays	Shotstack
Knowledge ingestion	Website crawl + GAS Google Sheet → Drive folder

Connectors (per tenant, BYO): each agency/user connects their own ElevenLabs, Higgsfield, HeyGen, and Magnific accounts. GAS connects GAS's for Iteration 1. See [ux-flow.md](#) .

Tool roles (your existing stack, assigned deliberately — cost-sensitive, no quality compromise):

- **Claude Code in VS Code / Codespaces** — the build environment. Read the Claude Code skills for HeyGen, Higgsfield, ElevenLabs, Magnific (MCP) and Vercel storage to wire current, non-deprecated patterns.
- **Magnific** — realism only, and *selectively*: Skin-Enhancer/Precision on the influencer's identity reference frames and per-video hero frames (~\$0.08–0.16/image), not every frame. Low creativity so it adds pores without changing the face.
- **Google Drive** — knowledge ingestion (GAS Sheet → folder) and optional final delivery into a client Drive folder.
- **Gmail** — notifications and delivery: "ready to approve", "render complete", client hand-off emails.
- **Zapier** — *edge glue only* (notifications, hand-offs, future publishing). The core render pipeline stays in **Inngest** — durable

and in-code; do not rebuild the pipeline in Zapier.

- **Canva** — optional, narrow: brand-kit/thumbnail/static-post generation. Keep it OUT of the video render path (Shotstack owns in-video overlays) to avoid overlap.

Output formats: 45s, with `aspect_ratio` per platform — 9:16 (Reels/TikTok/Shorts), 1:1 or 16:9 (LinkedIn). Composed for the format, not cropped after.

Irreducibly external: Claude, ElevenLabs, HeyGen, Higgsfield, Magnific, Shotstack.

Stack audit (June 2026) — why we use

Higgsfield's *models*, not its *Supercomputer*:

Higgsfield Supercomputer (launched May 2026) is an agentic layer that plans a brief, routes

across frontier models, and ships assets — essentially a competing orchestrator. We do NOT

build on it as the engine, because: (1) you can't white-label/productize someone else's agent as your own SaaS; (2) its credit model has

dynamic throttling and opaque cost — the opposite of the hard, per-stage cost gate this agency needs; (3) it can't host our client-brain RAG, consent/compliance, per-member dashboard, or our UX. We keep **our** Inngest orchestration (owned IP, deterministic, auditable, cost-gated) and call Higgsfield's best-of-breed models (Soul 2.0, Seedance 2.0, Cinema Studio) via API. It's useful internally as an ideation scratchpad, and it validates that our orchestration pattern is the right one.

Model policy (cost-accurate): the producer LLM is ~3% of a ~\$3.50 video — Opus vs Sonnet is ~\$0.10/video. So: **Sonnet 4.6 default**, Opus 4.8 (effort=high) as an opt-in premium pass, Haiku 4.5 for ingestion, prompt caching always on. Control cost on render dials (HeyGen avatar length, Higgsfield tier, resolution, Magnific frame count), not the model.

See `cost-controls.md` for the worked numbers.

Deprecated — never use: `@vercel/postgres` ,

@vercel/kv . Use

@neondatabase/serverless .

The product flow (see [ux-flow.md](#) for the full UX)

Setup (occasional, reused):

- **Connect Tools** — ElevenLabs / Higgsfield / HeyGen credentials per tenant.
- **Build an Influencer** (synthetic, from the company/brain) **or Build Me** (digital twin: upload photo + connect ElevenLabs voice or upload a 45s voice file). Reusable assets.
- **Brains** — connect website + GAS Sheet → Drive knowledge.

Produce (the repeatable journey, on the "build spine"):

Brief → Script & Scenes (producer co-pilot) → Review & Approve (cost gate) → A-Roll → Build B-Rolls → Add Sounds (optional music / optional SFX) → Stitch Video → QA → Preview & Download.

The design docs

- `docs/ux-flow.md` — connectors, the produce journey, the producer co-pilot, glass-box overlays, world-class UI/UX direction + tokens.
- `docs/influencer-builder.md` — the two identity modes, the embedded hyper-realism master prompts, voice options.
- `docs/producer-prompt.md` — the producer co-pilot's system prompt (the core quality component).
- `docs/compliance.md` — POPIA/GDPR express consent, the compliance register tab, retention/erasure.
- `docs/architecture.md` — Neon schema, tenancy, isolation, productization seams.
- `docs/brains.md` — knowledge ingestion (website + GAS Sheet → Drive), RAG. (Learning loop = Phase 2.)
- `docs/production-pipeline.md` — the 9-stage video factory, 45s structure, draft-then-final, QA gate.

- `docs/audio-direction.md` — dynamic music; optional per-scene SFX/ambient.
- `docs/cost-controls.md` — rate card, estimation, budget gate, dashboard.

Build phases (Iteration 1 unless noted)

1. **Foundation** — Next.js/Vercel; Neon + pgvector + Blob; Auth.js; schema; brain dropdown.
2. **Connect Tools** — per-tenant credential vault for ElevenLabs/Higgsfield/HeyGen/Magnific.
3. **Influencer Builder + consent** — synthetic + digital-twin modes; hyper-realism prompts; Magnific Skin-Enhancer on reference frames; Soul/voice binding. **Express POPIA/GDPR consent gate at every photo/voice upload, with timestamp, and a Compliance register tab.**
4. **Brains** — website crawl + GAS Sheet → Drive sync; profile builder. Pass the isolation test.

5. **Script + producer + estimate** — producer co-pilot (`producer-prompt.md`); rate-card estimate.
6. **Cost gate + dashboard** — `rate_card` , `budgets` , runs; hard pre-flight gate; per-member dashboard.
7. **Production pipeline** — a-roll (HeyGen), b-roll (Higgsfield) + selective Magnific hero-frame pass, optional audio, draft assemble (Shotstack); Gmail notifications.
8. **Finalise + QA + deliver** — draft gate, full render, glass-box overlays, QA gate, Blob download.
9. **Iteration 2 (market + intelligence):**
performance/learning loop, Meta feedback, ad-loading, white-label.

Quality gates: the producer self-verifies the script before emitting; the QA gate verifies the finished cut on a tiered-severity model before delivery. Never stop at a first draft.