

Influencer Builder — Believable AI Humans

An influencer is a **reusable identity** built once and used across many videos: a Higgsfield **Soul** (consistent face/identity), a voice, and a look reference set. The entire point is that it reads as a real person — so the realism prompt library below is embedded and applied automatically; the user never has to know how to prompt for skin pores.

Two modes. Both produce the same reusable record on `clients` (or a `influencers` table if you want several per brain):

`higgsfield_soul_id`, `voice_id`, look refs, brand fit.

Mode A — Build an Influencer (synthetic, from the brain)

For a persona that represents the company,
not a real person.

1. **Derive a brief** from the active brain:
audience, market, tone → a short persona
spec
(age range, vibe, wardrobe, setting) the
producer co-pilot drafts and the user
tweaks.
2. **Generate reference frames** with Higgsfield
using the **hyper-realism master prompt**
(below) + the persona spec. Generate a
small set; the user picks the 3–5 most
believable.
3. **Train Soul 2.0** on those frames →
`higgsfield_soul_id`. This locks identity
across every
future scene and video (no drift). Soul 2.0 is
the current identity model; pair with
Seedance 2.0 / DoP + Cinema Studio for
the b-roll motion.
4. **Assign a voice** — pick or clone an
ElevenLabs voice that fits the persona →
`voice_id`.

5. Save. The influencer is now selectable in the produce flow.

Mode B — Build Me (digital twin, authentic)

For the user's own likeness — the most authentic option.

1. **Upload photo(s)** of the person (clear, well-lit, a few angles). **Express consent gate fires here** — the upload is blocked until POPIA/GDPR consent is ticked and recorded with a timestamp (see `compliance.md`). Same gate applies to the voice upload in step 3.
2. **Create the avatar** — HeyGen photo/Avatar V from the still(s) for talking-head a-roll, and/or train a Higgsfield Soul from the uploaded photos for b-roll consistency.
3. **Voice — two authentic paths:**
 - **Connect ElevenLabs** and select an existing voice (e.g. a prior clone) → `voice_id`.

- **Upload a ~45-second voice file** → clone it in ElevenLabs → `voice_id`. This makes the influencer sound like the real person, which is the whole reason for this mode.

4. Save. "You" is now a reusable influencer.

Consent/rights: only allow likeness/voice upload for a person who has authorised it. Store a consent flag on the record. This matters more once productized to outside customers.

The hyper-realism master prompt (embedded; applied automatically)

Believability comes from texture and imperfection, and from killing the "AI sheen."

Apply

this to every identity generation and every b-roll frame featuring the influencer. Expose

only simple human controls (e.g. "setting", "wardrobe") – the realism scaffolding is hidden.

Positive core (always on):

photorealistic, natural skin with visible pores and fine vellus hair, subsurface scattering, subtle imperfections and natural asymmetry, realistic catchlights in the eyes, natural under-eye area, minimal or no makeup, soft directional key light with gentle falloff, shot on 85mm at f/2.0, shallow depth of field, neutral filmic color grade, relaxed candid expression, true-to-life proportions

Negative core (always on):

plastic or waxy skin, airbrushed, over-smoothed, beauty-filter, doll-like, CGI/3D-render look, uncanny symmetry, oversaturated, glossy plastic highlights, HDR halo, over-sharpened, excessive makeup, mannequin

Consistency rules:

- Always generate the influencer via their `soul_id` — never a fresh face per scene.
- Lock a seed per influencer for repeatability; vary only setting/wardrobe/angle per scene.
- Keep lensing and grade consistent across a video so scenes feel shot by one DP.

Magnific realism pass (the pore-level finish):

Run the chosen identity reference frames — and later, each video's hero frame(s) — through Magnific's **Skin Enhancer / Precision** endpoint at **low creativity**. This adds real skin texture, pores, and micro-detail that pushes past the AI sheen, *without* altering the identity (high creativity would change the face — don't). It's pay-per-image (~\$0.08–0.16), so apply it selectively to the frames that matter, not every generation. Enhance reference frames *before* training Soul so the locked identity is already pore-accurate.

Believability checklist (producer self-checks before accepting a frame):

skin reads as skin (texture, not wax) · eyes have natural catchlights and aren't glassy · expression isn't frozen · lighting has a real direction · no smeared hands/ears/teeth · identity matches the Soul. Regenerate any frame that fails — don't ship the uncanny one.

What the produce flow consumes

Downstream stages just reference the chosen influencer: a-roll drives the HeyGen avatar with the influencer's audio; b-roll generates scenes via the Soul; the voice stage uses `voice_id` (or the uploaded audio). Build once, reuse everywhere.
EOF